

QUORUM

GTM Context & Handover

Customer & lead pipeline, product signal, and working playbook — as of July 7, 2026

This document consolidates everything on the customer/lead GTM side to date: the product, the ICP, the pipeline, the operating rules that have been learned the hard way, the messaging scripts in active use, and a sample session debrief showing how sessions get turned into product and GTM decisions. It intentionally excludes the founding-team / hiring conversation — that is a separate, ongoing thread and isn't the subject of this doc.

1. Executive Summary

Quorum (product name: Mirror) is an AI-powered decision-support tool. A council of five distinct AI advisors responds to a real decision the user is sitting with — a career call, a business pivot, a major life choice. Launched early June 2026; this document reflects ~5 weeks in market.

Where things stand — July 7, 2026

- Revenue: ₹0. Paid subscribers: 0. First paying customer not yet secured — this is the single biggest open item.
- Sessions completed: 11. Mirror (paid) asks made: 3. Mirror access gifted to 1 strategic connector.
- Active pipeline: 10 people currently being worked. Parked leads: 13 — parked means on a structured, dated re-engagement window, not abandoned.
- Two organic advocates have emerged without being pitched — Viral Malwania and Gagan Dhir are using and sharing Quorum unprompted. This is the strongest signal so far.
- One inbound channel is live: an IIM-A alumni group post is up since June 30. It produced its first real inbound question July 5 (a comparison question, answered same-day). July 10 is the checkpoint for whether that inbound converts into an actual booked session.

The one governing principle behind everything below

Never explain the product before running it live. Conviction comes from experience, not description. The 30-minute live session on a real decision is the product itself — not a demo, not a sales step. Every script, every rule, and every piece of pipeline logic below is downstream of this.

2. Product Overview — What Quorum / Mirror Is

Mirror is the personal daily journaling / decision-reflection layer. Quorum is the council session built on top of it. When a user brings a real decision, five AI advisors respond, each from a different cognitive and philosophical angle:

- Elder — wisdom, long view, values, what will matter in 10 years
- Mirror — reflects what the user is already feeling and weighing, without judgment
- Scientist — data, evidence, probabilities, what the research says
- Strategist — options, trade-offs, sequencing, second-order effects
- Examiner — challenges assumptions, stress-tests the reasoning, finds the blind spot

Recently fixed: the Examiner now fires on every single session without exception. A rule engine (R1–R12) detects upstream blockers, irreversibility, missing information, value conflicts, time pressure, and multi-party dynamics. An emotional/gut question (E0) and a JTBD context-grounding question (C0) always fire too. This closes the "left-brain only" gap that early users flagged.

The one hard rule for running sessions

Do not explain the product before running it live. Do not pitch features. Do not describe the advisors in the abstract. Run a real decision, live, and let the person discover what happened. Every conversion so far has followed this pattern; every stall has come from explaining too early.

3. Ideal Customer Profile

3.1 Primary ICP

Dimension	Detail
Education	XLRI / IIM / equivalent tier-1 MBA or postgrad
Age	30–50. Peak earning years. Stakes are real.
Role	Senior professional, CXO, founder, operator. Owns decisions.
Situation	At a major crossroads — career shift, business pivot, life decision. Not a daily consumer.
Mindset	Introspective. Uses structured frameworks at work. Comfortable with reflective tools.
Income	₹25L+ per year. ₹999/month is negligible if the value lands.
Network	Warm introduction via XLRI alumni, IIM-A network, or second-degree founder connection.

3.2 Anti-ICP — do not pursue

- People looking for therapy or emotional support — Mirror is not that
- People in active crisis — not the right moment
- People who want to be told what to do — Mirror shows what you are already weighing, it does not prescribe

- Corporate HR or L&D buyers looking for a training tool — wrong channel, wrong pitch

4. Pricing

Tier	Rate	When to use	Notes
Founding Monthly	₹999/month	Default ask, first 20 paying users	Lock this now — price will rise
Standard Monthly	₹1,999/month	After founding cohort closes	Do not offer yet
Founding Annual	₹3,999/year	Only if user asks about annual	~40% saving vs. monthly founding
Gift access	No charge	Connectors & high-trust referrers only	Not a trial. Not for pipeline leads.

Never: drop below ₹999. Never offer free trials to pipeline leads. Gift access is reserved for strategic connectors. Never offer the standard rate before the founding cohort closes.

How to present the ask (verbatim)

“I am offering founding access at ₹999 a month — that rate locks in for as long as you stay. I am only extending this to people in this first wave. It takes about 30 seconds to set up.”

Never announce the price unprompted. Never negotiate on price. If someone pushes back on price, the real problem is conviction, not cost — the session didn't fully land.

5. Active Pipeline — as of July 7, 2026

5.1 Hot — active this week

Person	Last touch	Status	Next action
Girish B Hukkeri	Jun 27 (high-EQ response)	Personal financial crisis. Connector track fully paused.	Await. Do not mention Quorum. Respond as human first if he surfaces.
Gagan Dhir	Jul 6 (Day 12 nudge)	Daily WhatsApp running. Session 11 done Jun 24. Day-lock call still unconfirmed.	Day 13 nudge tonight 7 PM. No day-lock re-ask before ~Jul 9.
Viral Malwania	Jul 6 (Day 12 sent)	Organic advocate. Founder sharing. NOT a customer target.	Day 13 daily WhatsApp today 7 PM. No other contact.
Sujitesh Das	Jun 25 (3 messages)	Left-brain critique. Silence mode since Jun 25.	Window opens Jul 7–10 — milestone-only re-contact (Examiner fix).
Aditya Pittie	Jun 25 (email, no reply)	Email channel closed. LinkedIn is the right channel.	Short LinkedIn DM only.
Abhigyan Dutt	Jun 20 (replied)	Said he'll share a time slot. He owes the move.	Await.
Abhishek Agarwala	Jun 26 (LinkedIn sent)	First message sent. Awaiting.	Await reply.
Jay Ingle	Jun 9 (20 days)	Lowest-priority active lead.	Light touch sent Jun 29. Park if no reply by Jul 7.

5.2 Warm / in-pipeline — slower cadence

Person	Last touch	Status	Next action
Shivani	Jun 23 (check-in)	Mirror gifted. Session done. Friend track.	No product ping. Let her bring it up.
Dora Suri	Jun 27 (soft check-in)	Mirror gifted Jun 9. Replied, hasn't used it yet. Warm non-rejection.	Parked ~3 weeks. Re-engage ~Jul 18 with a genuine new hook only.
Sarthak	Jun 18 (reschedule)	Day 10+ silence. Still in pipeline.	Jul 8 light logistics touch. No product frame.
Puneet	Jun 7 (21 days)	21 days silent.	Jul 7 re-engagement — calibration-layer hook.
Yogesh	Never	Post-Sarthak referral only.	After Sarthak confirms.

5.3 Second-order — not yet contacted

Person	Connection	Status	Next action
Mohit Gupta	IndiaMART Head of Products. Old NSIT contact.	Not yet contacted this wave. Good ICP fit.	Week of Jul 7 at earliest. First-paying-user hook preferred.

Person	Connection	Status	Next action
Vaibhav	Ex-EXL, IIM-A, now Kearney Investments.	Not yet contacted this wave. Good ICP fit.	Week of Jul 7 at earliest. First-paying-user hook preferred.

6. Re-engagement Calendar — No Lead Closes Permanently

Every parked or closed lead has a structured re-engagement window. The window is based on how the lead disengaged: the more abruptly they went quiet, the longer the gap before contact. Re-engagement always uses a new angle — never the same ask repeated.

Spacing logic

Gap	When it applies
20 days	Still nominally active but silent — in pipeline, last touch under 3 weeks.
30 days	Went quiet after initial interest — 1–2 messages, faded, no hostility.
40 days	Complete ghost — multiple messages sent, zero engagement, sessions missed.
90 days	Bad faith or fundamentally misaligned. Only if a strategic angle materialises.

Schedule

Lead	Last touch	Re-engage date	New angle
Surbhi Allagh	Jun 20 / nudged Jun 29	Aug 3	30-day gap — genuine new hook, not a repeat of the reactivation ask.
Jay Ingle	Jun 9	Jun 29	“Didn't want this to go fully cold. Offer still stands.”
Aditya Pittie	Jun 25 (email)	~Jun 30 / Jul 1	LinkedIn DM only. Short. One hook.
Navneet Bali	Jul 3 (5th same-ask msg)	Jul 31	New angle only — over-contacted at 5 msgs/12 days.
Sarthak	Jun 18	Jul 8	Logistics only. No product frame.
Puneet	Jun 7	Jul 7	Calibration layer — tracks confidence before/after outcomes.
Kovil	Jun 10	Jul 10	Decision graph + memory architecture.
Yugal	Jun 12	Jul 12	“First paying user. Made me think of you.” Judgment record frame.
Achint Sachdeva	Jun 14	Jul 14	Decision continuity — sessions link for layered venture build.
Sumit	Jun 15	Jul 15	XLRI + milestone: sessions ran through it.
Chirag	~May 20	~Jul 15	Clean break. No reference to old conversation.
Dhruv Khandelwal	Jun 16	Jul 26	“Product revamped. First paid user.” No session ask.
Manmeet Singh	Jun 20	Aug 3	30-day gap. Genuine milestone hook only, not a status check.

Lead	Last touch	Re-engage date	New angle
Rakesh Santhapur	Jun 20	Aug 3	Same as Manmeet.
Mohit Gupta	Never this wave	Week of Jul 7+	First-paying-user hook. NSIT/IIM angle.
Vaibhav	Never this wave	Week of Jul 7+	First-paying-user hook.
Dezerv	May/Jun (closed)	~Sep 2026	Only if a genuine strategic angle materialises. Do not chase.
Dora Suri	Jun 27	~Jul 18	Genuine new hook only. Do not mention referral.

Three rules for every re-engagement

- Never repeat the original ask. If it didn't work once, change what you're asking for, not the volume or tone.
- Never reference the old silence. Don't open with "following up again." Start clean.
- Never send without a genuine new hook. A product milestone, first paying user, or new use case. A weak hook means wait for one.

7. Daily WhatsApp System

Established June 25, 2026. Currently running for Gagan Dhir and Viral Malwania. Purpose: maintain a daily touchpoint with high-engagement users without creating social debt or requiring replies.

DB signal (Jun 28): neither Gagan nor Viral is running sessions daily. This is expected — real decisions don't arise every day — and does not invalidate the system. The daily message maintains top-of-mind presence so Quorum is the first move when a decision does arise. Do not adjust format, frequency, or content based on this signal.

Format rules (locked)

- Warm personal opener — acknowledge the person, not just the product
- One light reference to Mirror — a reminder, not a check-in or progress question
- Explicit "no reply needed" — every message, non-negotiable
- Never a question that requires an answer
- Never name-drop other users
- Vary the phrasing slightly each day — same structure, different words

Templates

Gagan (formal/warm register): "Good morning Gagan sir 🙏 Dropping your daily Mirror reminder — no need to reply to this, just there whenever something comes up today."

Viral (casual/energetic register): "Hey Viral 😊 Daily drop — Mirror's there whenever. No reply needed, just keeping the habit alive!"

Periodic checkpoint calls (e.g., the Jun 30 Gagan call)

Two objectives every time: (1) usage diagnostic — what has he actually run through it, what decisions, what was useful; (2) referral unlock — specific names from his network who are at a real crossroads. Close with a direct ask for 2–3 specific names rather than a generic “know anyone?”

8. Connectors & Referral Sources

Girish B Hukkeri

Field	Detail
Role	Founder — i2u.ai / embryoFund ecosystem
Value	Wide network across founders, CXOs, operators — ideal ICP overlap
Status	Personal crisis disclosed Jun 26 (loan repayments, VC milestones pending). High-EQ response sent Jun 27. Full stop on Quorum contact.
Next step	Respond as a human being if he surfaces. No agenda. Connector conversation resumes only on his terms.

Vishal Arora (Mastercard / IIM-A alumnus, based in Canada)

Field	Detail
Relationship	IIM-A connection. Ran an exploratory session earlier in the playbook period. Fed back on HNI benchmarking and personality-weighted outputs — both features now shipped.
IIM-A group post	Live since Jun 30. First inbound received Jul 5 — a group member asked how Quorum compares to a competing product; answered same-day. Has not yet converted into a booked session.
Status	Active monitor state. Reply within 1–2 hours of any further inbound. Jul 10 checkpoint measures whether the inbound converts into a booked session — if the thread goes quiet again with nothing booked, move to a personal, direct message rather than another public post.
Inbound protocol	Speed matters. Don't pitch, don't send the link first — invite instead: show, don't explain. If asked for the link before a session, hold it and offer to run it live instead. At the close of any inbound session: make the paid ask once, then silence.
Do not	Send another nudge about the post itself — it's live, nothing left to nudge. Never send the link before a live session with an inbound contact.

Raghvendra — friend-support track

Field	Detail
Relationship	Close friend. Ran a session very early, treated as feedback rather than a conversion session.
Network value	IIM-A + BCG-adjacent connections, directly overlaps ICP. Referral value retained but not currently pursued.
Status	Reclassified to friend-support track. Quorum contact fully paused after an in-person meeting surfaced substantive, firsthand skepticism about AI solving judgment-level problems — a grounded, personal-failure-based objection, not casual doubt. It does not respond to reassurance or feature-selling.
Next step	No Quorum content unless it surfaces naturally from him. Respond as a friend first, no strings, no scheduled re-engagement date. Referral ask is shelved indefinitely, not abandoned.

9. Key Lead Briefs

Navneet Bali — parked (over-contacted)

Session done Jun 11. Mirror ask made — he said “let me think, in person works better.” Between Jun 21 and Jul 3, five consecutive messages went out on the same in-person scheduling ask (Jun 21, 23, 25, 29, Jul 3), all soft “circling back / no pressure” variants. Zero replies across all five over 12 days — this exceeds the three-message cap. Each message read as low-stakes individually, but the cap counts messages sent, not tone.

Status: parked as of Jul 4. No contact until July 31. The reset message must carry a genuinely new angle — never a status check on the same ask.

- Format when eventually confirmed: 30–40 minutes in person. Show, don't explain. Payment set up in the room.
- Track 1 only: Mirror at ₹999/month founding rate. Close this before anything else.
- Track 2 (org/team use) only after Track 1 closes — never mention it before then.
- Do not mention the ₹3,999 annual rate unless he asks. Do not explain the product's ontology or council design.

Gagan Dhir — daily WhatsApp + periodic checkpoint calls

Session 11 done Jun 24. Self-described Quorum ambassador. Wants to use it in client decision contexts (org/CXO embedding). Running on the daily cadence — DB shows no daily sessions, which is expected (see §7).

Product update to share when a call happens: the Examiner now always fires (his own backlog item) — rule engine R1–R12 is live, and the emotional/gut question (E0) always fires too. He can now recommend it to a client with confidence.

Day-lock ask vs. daily nudge: permanently separate tracks. The daily “no reply needed” nudge and the day-lock scheduling ask must never be merged into one message. The day-lock re-ask is a single, standalone, daytime message, sent at most once a week. If he names a day, lock it in one line and say nothing more. If he deflects or stays silent, drop it — don't chase — and let him come back to it on his own.

Sujitesh Das — silence mode

Silence mode was in effect until the Jul 7–10 window. Do not send any message outside a genuine trigger. Next contact must be tied to a product milestone or a specific insight in his world — “checking in” is not a valid reason. Three messages went out Jun 25 (one good reply, two unnecessary follow-ups) — do not repeat that pattern.

Raghvendra — friend-support track

First outreach (one-hook version) went out Jul 1 and he replied, leading to an in-person meeting. That meeting surfaced substantive, firsthand skepticism about AI solving judgment-level problems, rooted in his own attempt to automate judgment-heavy work with AI and watching it fail. This is a distinct objection type — grounded, personal-failure skepticism, not casual doubt — and it does not respond to reassurance, more features, or a reframed pitch.

Status: Quorum contact fully paused. Moved off the active pipeline entirely. Referral/network value is retained but shelved, not pursued.

Going forward: respond as a friend first, no strings attached. No Quorum content, no re-engagement date, no scheduled next touch. If Quorum comes up naturally in his own words, that's the only valid re-entry point — don't manufacture one.

10. Messaging Scripts Library

These are the working scripts in active use across the pipeline. Each exists because an earlier, more generic version underperformed — treat the specific wording as deliberate, not decorative.

Script 1 — First outreach — warm connection

“Hi [Name] — [shared connection] mentioned you and I wanted to reach out directly. I've been building something that I think you might find genuinely useful — an AI council for major personal decisions. Not a productivity tool. Something you sit with. Would you be open to a 30-minute conversation? I'd rather show you than explain it.”

Do not: send a link, explain the product, mention pricing.

Script 2 — Session invitation — after warm intro

“[Name] — I'd love to run a live session with you rather than just describe what it does. Bring something you're actually sitting with — a career call, a business decision, something personal. Doesn't have to be resolved or even well-formed. We do 30 minutes. If it doesn't land for you, no harm done.”

Do not: explain the council in advance, offer examples, over-explain.

Script 3 — Post-session follow-up (within 24 hours)

“[Name] — how is it sitting? Anything that stayed with you from yesterday?”

Do not: ask whether they want to subscribe, explain further, follow up twice without a reply.

Script 4 — Second follow-up (if no reply after Script 3)

“[Name] — no rush, just leaving the door open. If something comes up that you want to run through it, I'm around.”

Do not: mention pricing, make this a third follow-up.

Script 5 — Mirror paid ask — direct, at close of session

“Before we close — I want to offer you founding access to Mirror. It's ₹999 a month. That rate locks in for as long as you stay. I'm only extending this to the first wave of people. Takes about 30 seconds to set up. Are you in?”
Then: silence.

Do not: offer before the session is complete, negotiate, mention the standard rate.

Script 6 — Reactivation — after long silence

“[Name] — dropping back in. Something happened with the product since we last spoke that made me think of you. [One specific thing]. If the timing is better now, I'm happy to pick this up.”

Do not: reference the silence, reference the last message, ask what happened.

Script 13 — Paid ask — full version, in-person

“I want to offer you something before we wrap up. Mirror — the daily layer — is now available for founding members. ₹999 a month. That's the rate I'm holding for the first 20 people. It locks in for as long as you stay. Takes 30 seconds to set up right now if you want to do it here.” Then stop talking.

Do not: fill the silence, explain further, mention ₹3,999 unless asked.

Script 15 — Gifted-user soft check-in

“Hey [Name] — how's Mirror sitting? Any decisions you've run through it?” Nothing more. Referral angle only after a genuine positive reaction.

Script 16 — Left-brain critique response

"I hear you — and I want to push back slightly. The Mirror advisor and Elder in the council are specifically designed to surface the social and emotional weight of a decision, not just the functional logic. If the session felt left-brain heavy, that's a fair critique of how the session ran, not a verdict on the product. What's one thing you're actually sitting with right now?" Can now add: "The Examiner gap has been fixed — it fires every session now, with a question that goes straight to the gut and emotional layer."

Do not: send a second message before they answer, reference how others are using it, get defensive.

Script 17 — Daily WhatsApp reminder — two templates

Gagan (formal): "Good morning Gagan sir 🙏 Dropping your daily Mirror reminder — no need to reply to this, just there whenever something comes up today." Viral (casual): "Hey Viral 😊 Daily drop — Mirror's there whenever. No reply needed, just keeping the habit alive!"

Script 18 — Long-delayed first outreach

"[Name] — I should have looped back sooner. We connected earlier this year. Since then I've actually run the product live with 11 people — real decisions, not demos. I'd value your read on it. 15 minutes this week if the timing works?"

Do not: reference the gap or silence, apologise, explain the product in the outreach.

Script 20 — Product-evolution re-engagement for parked leads

Adapt one specific hook per person — e.g. calibration layer (Puneet), decision graph + memory architecture (Kovil), decision continuity across sessions (Achint). Lead with the specific thing that changed, not a feature list.

Do not: send a feature list, lead with mechanism before outcome, use more than one variant on the same person.

Script 21 — One-hook daily nudge / day-lock separation

Daily nudge (no reply needed, varied daily) and day-lock scheduling ask (standalone, daytime, sent separately and less often) are two permanently separate message types for anyone on a daily cadence. If he names a day, reply only: "Perfect — [day/time] works. Will connect then." Nothing more.

Do not: fold a question into a "no reply needed" message, send more than one day-lock re-ask per week without a new signal.

Script 19 — Close-friend re-engagement (one hook only)

"Hi bro — how's [the thing they mentioned] panning out? Also, remember you flagged [the specific gap] — that's fixed now. If there's something top of mind, let's run it again? 20 mins, whenever available?" Casual friend register, one clear ask, no feature stacking.

Do not: stack multiple product hooks into one message, send a second message before they reply, go longer than 3 short lines.

Script 22 — Competitor / comparison inbound response

Structure: (1) one sentence on what the comparison product actually does, in its own terms; (2) one sentence on what Quorum does that's a different category entirely, not a better version of the same thing; (3) close by inviting them to experience it live rather than continuing the debate in-thread.

Do not: answer with a feature comparison table, debate the other product's merits, let the thread run past one reply without an invite.

11. Product State — What's Been Built

Use this when crafting re-engagement or sales messages. Not a feature list to send verbatim — translate each item into the specific language of the person you're messaging.

Feature	Detail
Examiner always fires	Rule engine (R1–R12) detects upstream blockers, irreversibility, missing info, value conflicts, time pressure, multi-party dynamics. E0 (emotional/gut) and C0 (JTBD context) fire every session. Resolves the earlier left-brain critique — safe to tell anyone it's fixed.
End-to-end encryption (AES-256-GCM)	All decision text, context, Examiner responses, and outcomes are field-encrypted before reaching the database. Anyone with direct DB access sees only opaque strings. Relevant for senior professionals with real, sensitive decisions.
Decision Graph	After enough sessions, users get an interactive visual map of how their decisions connect — structural similarities, shared biases, contradictions, user-asserted links. Force-directed, click-through to individual records.
Outcome Tracking + Calibration Fingerprint	Users log what they actually decided and whether the council helped. The product compares pre-decision confidence to retrospective confidence after the outcome — building a bias fingerprint over time: where judgment holds, where it doesn't.
Decision Continuity	When a user revisits a decision, the council reasons from the parent session's synthesis, Examiner answers, commitments made, and any logged outcome. Memory that compounds across sessions.
Monthly Judgment Review	Rolling 30-day snapshot: decisions total, loops closed, patterns confirmed, rule recall applied. A structured accountability layer.
Push Notifications	30 nudge variants across 7 themes, personalised with session count and bias label, rotated so nothing repeats within a month.

12. Product Backlog — Surfaced from Real Sessions

Items below came out of a live session (Gagan, Session 11). They are product items, not GTM blockers — flagged here so anyone on the GTM side knows what's open vs. resolved when a prospect raises something similar.

Item	Status
Context as a first-class input (mood, fear state, archetype active shifts the right answer)	Open
Samskara / context layer — surface which archetype is most active in a user's framing	Open
Calibration/profiling made visible at onboarding, not buried	Open
Risk-aversion bias surfaced explicitly to the user	Partially resolved — Bias Fingerprint now live
Examiner should fire every time	Resolved — rule engine R1–R12 live
Org/CXO embedding risk — one CXO's biases embedded in an org tool means the org inherits them	Open — relevant to enterprise conversations
HNI benchmarking (needs a larger decision corpus to be accurate)	Partially built

13. Open Flags

- Session-count reconciliation: an earlier note flagged one additional completed session reflected in the running count but never individually logged. This has never been fully reconciled — the count of 11 sessions assumes it's resolved. Worth double-checking the session log directly.
- Everything else previously flagged (IIM-A post status, Navneet's over-contact correction, Raghvendra's reclassification, Vishal's checkpoint logic) has been resolved and is reflected accurately in the sections above.

14. Locked Decisions — Do Not Re-litigate Without New Information

- Pricing floor is ₹999/month. Never drop below it.
- Never offer free trials to pipeline leads. Gift access is for connectors only.
- Never share the Mirror link before completing a session.
- Maximum three follow-up messages per cycle to any single lead. Then park and reset.
- The paid ask is made live, in session, at the close. Never in a follow-up message.
- ICP is XLRI/IIM-tier professionals at a real crossroads, 30–50. Do not chase outside it.
- The ₹3,999 annual rate is mentioned only if the prospect asks. Never offered proactively.
- The public LinkedIn post is held until the first paying customer is secured. No exceptions.
- Dora is a gifted connector, not a prospect. The referral angle is paused until she shares genuine organic feedback.
- Navneet's pitch follows Track 1 (personal) before Track 2 (org use) — never in the same conversation.
- Viral is not a paid-customer target. He's an organic advocate — no paid ask to Viral.
- Girish is a connector, not a customer. The session objective is specific referral names, not his own subscription. Currently: full stop, personal crisis.
- The session runs on the prospect's real decision — never a demo scenario or a prepared hypothetical.
- Shivani is on the friend track. No Mirror push — let her bring it up.
- Dezerv: closed. No further contact without a genuinely different strategic angle. 90-day minimum.
- Dhruv: parked at a 40-day window. Never ask for a session as the first ask on re-engagement.
- No more than two messages to the same person in a single day. One is ideal.
- Session count must be tracked accurately — reconcile any gap promptly.
- Re-engagement always uses a new angle. The same ask, repeated, is harassment, not re-engagement.
- IIM-A post via Vishal: do not send another nudge before a real milestone — the post itself needs no more nudging.
- Sujitesh: silence mode until the window opens; next contact must be tied to a product milestone, never a check-in.
- Viral: organic advocate track only, permanently. Referral value exceeds subscription value.
- Daily WhatsApp format is fixed: warm opener + light reminder + explicit “no reply needed.” No questions, no comparisons to other users.
- Aditya Pittie: LinkedIn is the primary channel going forward — email has been tried and didn't work.
- Raghvendra: first message trimmed to one hook, not a stacked multi-feature draft. No second message until he replies.
- Vishal's IIM-A post: monitor for inbound, reply within 1–2 hours, don't nudge the post itself again.
- Mohit Gupta and Vaibhav: no contact until week of Jul 7 at earliest, and only with a first-paying-user hook or a strong named-connector reference.
- The day-lock/scheduling ask and the recurring daily nudge are permanently separate tracks for anyone on both — never merge a question into a “no reply needed” message.
- One hook per message. Never stack multiple product features, questions, or callbacks into a single outreach message.
- Navneet Bali is parked as of Jul 4 — over-contacted at 5 messages/12 days. No contact until Jul 31. The reset message must carry a genuinely new angle.
- Vishal's IIM-A post has a hard reassessment checkpoint of Jul 10 — measuring whether inbound activity has converted into a booked session.
- Raghvendra is reclassified to friend-support track. Quorum contact fully paused. Referral value retained but not pursued. No re-engagement date — only re-enters if he raises it himself.

- Comparison/competitor questions surfaced through a connector's network are treated as inbound — same 1-2 hour reply window, answered with one sharp category distinction rather than a feature-by-feature rebuttal, closed with an invitation to a live session.

15. Key Learnings

- Run it live, never explain it first. The product creates conviction through experience, not description.
- Lead with their decision, not the product. Senior users engage faster when the session is about something real, not a showcase.
- The session is the product. The 30-minute live council run is not a demo or a sales step — treat it as the product itself.
- Silence after a strong message is not rejection. Especially for senior, introspective people — the thinking takes longer than the typing.
- The paid ask belongs at the end of the session, not in a follow-up or a separate call.
- Never make a financial ask in a follow-up message — a follow-up ask is easier to decline than one made live.
- The right ICP person knows immediately. If they lean in and still don't convert, the problem is close mechanics, not conviction.
- Group posts are weaker than direct outreach. Direct warm messages convert 3–5x better than group posts, even in high-trust alumni groups.
- Referrals come from session experience, not awareness. People who've heard about it don't refer; people who've run a decision through it do.
- Never send a fourth follow-up in a single cycle. Three messages maximum before parking and resetting the window.
- Reactivation works best with a genuine new hook. “Following up” is not a hook — a milestone, a first paying user, or a named use case is.
- Price resistance is conviction resistance. If someone pushes back on ₹999, the session didn't fully land — don't negotiate.
- Never share the link before the session. The session creates the context; the link without it creates a weak first impression.
- Unprompted sharing is the strongest signal. “I will share this” is weak; sharing without being asked is the strongest possible signal.
- The “Mirror” frame outperforms the “AI council” frame. Lead with the output (“shows you what you're already weighing”), not the mechanism.
- Staged commitments are not nos. “Let me think” or “come back in a month” has a different trajectory than “not interested” — treat them differently.
- Gift access is for connectors, not prospects. Gifting removes the friction that makes conversion meaningful.
- Change the entry point, not the volume. If an ask hasn't worked, use a different angle — not a louder or softer version of the same one.
- Advocates don't start as enthusiasts. They emerge after sustained low-pressure presence, not during the pitch.
- Left-brain partial-surface critique: when someone only engages with part of the product, they may judge the whole from a partial experience. Address the gap in their experience, not the critique of the product itself.
- Organic-advocate arc — don't convert back. Once someone pivots from prospect to organic advocate unasked, don't re-pitch them as a customer; the referral channel is worth more.
- Three-messages-in-a-day trap. When a strong message ends on an open question, the only valid next message is their answer — sending more before they reply communicates anxiety and undercuts the close.
- Real decisions don't arise daily, and that's normal. The daily WhatsApp maintains presence — it isn't meant to drive daily session usage.

- Product evolution is a re-engagement hook only when specific and honest. “The thing you told me was missing is now built” works; “the product has evolved” doesn't.
- One hook per message. Stacking a personal question with two or three product features reads as a feature list, not a friend texting — no matter how warm the wording.
- Grounded personal-failure skepticism is a distinct objection type. When doubt comes from someone's own firsthand, specific failed attempt at something similar, it's not generic skepticism and doesn't respond to reassurance — park it, don't re-pitch until they raise it themselves.
- Low-key phrasing doesn't exempt a message from the follow-up cap. Five consecutive “just checking in” messages are still five follow-ups — audit cumulative count per ask cycle, not per-message tone.
- Comparison questions relayed through a connector are inbound signal, not idle curiosity — someone in that network is close enough to considering it to be checking against an alternative. Reply fast, land one sharp category distinction, invite them to a live session rather than debating in-thread.

16. This Week's Plan

16.1 Today — July 7 (Tuesday)

- Gagan — Day 13 nudge only. No day-lock re-ask (next valid window ~Jul 9).
- Viral — Day 13 nudge.
- Navneet — no message. Parked until Jul 31.
- Raghvendra — no Quorum contact. Continue waiting for him to raise it himself.
- Vishal — monitor the thread. Reply within 1–2 hours if it moves. No proactive push. Checkpoint Jul 10.
- Sujitesh — contact window opens today, milestone-only. The Examiner-fix hook is genuine and specific — send it today, one message only.
- Jay — checkpoint day. Light touch Jun 29, no reply since. Park him today per the standing rule. 30-day reset gap applies (~Aug 6).
- Manmeet / Rakesh / Surbhi — no action needed. Parked until ~Aug 3.

16.2 Today's messages (for tone/register reference)

Gagan — Day 13 daily WhatsApp (7 PM): “Gagan sir — Day 13 🙏 Just a small nudge, no need to reply — the space is there whenever you want it.”

Viral — Day 13 daily WhatsApp (7 PM): “Hey Viral — Day 13 🙌 No reply needed, just flagging I'm still here whenever it's useful.”

Sujitesh — milestone-hook re-engagement (one hook only): “Sujitesh — quick one. You'd flagged the sessions feeling left-brain heavy, Examiner not always firing or asking the right follow-ups. That's actually fixed now — it fires every session, goes straight for the emotional/gut layer too. Wanted you to know before anything else. If something's live for you, happy to run it again — 20 mins.”

Jay — park message (light, final touch before parking): “Jay — didn't want this to go fully cold, but I'll leave it here for now. Door's open whenever the timing's right.”

16.3 Week ahead

Date	What
Jul 6 (done)	Gagan + Viral Day 12 sent. Navneet/Raghvendra: no contact. Vishal: monitored, no change.
Jul 7 (today)	Gagan + Viral Day 13. Sujitesh re-engaged (Examiner-fix hook). Jay parked.
Jul 8	Sarthak light touch (logistics only, no product frame).
~Jul 9 onward	Earliest a fresh Gagan day-lock re-ask is appropriate, if still no reply.
Jul 10	Vishal checkpoint: has the IIM-A-thread engagement converted into a booked session? If not, move to a personal, direct message rather than another public post.
Jul 31	Navneet reset — re-engage only with a new angle.
Aug 3	Manmeet, Rakesh, Surbhi reset window.
Aug 6	Jay reset window.

17. Metrics Snapshot — July 7, 2026

Metric	Value
Revenue	₹0
Paid subscribers	0
Mirror access gifted	1 (Dora)
Sessions completed	11
Mirror (paid) asks made	3
Inbound questions answered	1 (IIM-A group, comparison question, Jul 5)
Active advocates	2 (Viral, Gagan)
Daily WhatsApp cadence	Gagan Day 13 · Viral Day 13
Vishal IIM-A post	First inbound Jul 5. No session booked yet. Checkpoint Jul 10.
Gagan call	No day locked. No re-ask until ~Jul 9.
Navneet	Parked until Jul 31.
Manmeet / Rakesh / Surbhi	Parked until ~Aug 3.
Jay	Parking today. Reset ~Aug 6.
Raghvendra	Friend-support track. Quorum paused.

18. Special Notes — Carry Forward

Vishal

No movement since the Jul 5 inbound. No session booked yet. The Jul 10 checkpoint measures conversion, not just presence of inbound — if the thread stays quiet with nothing booked, the next move is a personal message to Vishal directly, not another post.

Navneet

Parked since Jul 4. No further messages until Jul 31. When that window opens, the message needs an actual new angle, not “checking in.”

Gagan

Daily nudge continues unchanged. Day-lock re-ask window still ~Jul 9. Silence isn't a no.

Raghvendra

Friend-support track. No Quorum content unless it surfaces naturally in his own words.

Sujitesh

Window opened Jul 7 after the Jun 25 silence-mode period. The hook is specific and genuine — the exact thing he flagged is now fixed — so this qualifies as a real milestone hook, not a check-in. One message only; if he replies, engage naturally; if not, back to silence.

19. Open Items to Track Going Forward

- Vishal thread outcome — if the group replies further or a session gets booked before Jul 10, log it as a milestone and re-evaluate the checkpoint framing.
- Session-count reconciliation (the “session 10” discrepancy) — still unreconciled, worth a direct check of the session log.
- First paying customer — this is the single most important unresolved item. Everything in the pipeline is ultimately in service of this.

20. Sample Session Debrief — Dora Suri (May 23)

Included as a worked example of how a live session turns into product feedback and GTM decisions — this is the level of detail and the action/park discipline every session should get.

20.1 How the session went

A genuinely good first live session. Dora engaged with the product properly — ran the council herself, pushed back, had a real moment with the Examiner-style Risk Advisor, and started thinking like a product advisor without being asked.

- Worked — real moment landed: the Risk Advisor's caution against raising salaries clicked out loud. That's the product doing its job.
- Worked — unprompted product thinking: she immediately named a CEO client who makes high-frequency, high-stakes decisions as a use case, without being prompted. Strong product-market signal, not just feedback.
- Worked — sharp trade-off framing insight: “Each option comes with trade-offs — reputation, people, financial — and if Quorum shows those rather than picking a winner, the user stays in the driver's seat.” Worth sitting with seriously.
- Watch — her initial decision framing was upstream of the actual decision. It got caught and reframed live, but the product should probably flag this before the council runs, not after — the manual override is a product gap.
- Watch — the explanation of how outcome-tracking works before any outcomes are logged didn't fully land. The answer was correct but needs a cleaner one-liner.
- Watch — unclear whether the referral ask happened; understandable given a live session with product onboarding running simultaneously, but worth protecting in future sessions.

20.2 Product feedback — action vs. park

Five distinct pieces of feedback came out of this session. Three are sharp and actionable now; two are directional and need more data before acting.

Action now

- Trade-offs per option: add a trade-off layer to synthesis or per-advisor output (reputation / people / financial). Keeps the user in the driver's seat without fighting the product's existing lean.
- Delete-decision button: re-analysing a decision currently creates duplicates. Users should be able to delete copies and track outcomes only on the original — straightforward UX fix, meaningful for profiling clarity.
- Sub-text formatting: a piece of context copy currently reads as an instruction rather than context — italicise or visually demote it. Quick fix.

Park for now

- Persona/role labels: confusion cleared once explained live — the fix is a tooltip or one-line framing on the advisor card, not new labels or a framework rebuild.
- Back button after saving a decision record: valid UX point, low priority unless users hit it repeatedly.
- Upstream-dependency framing: a product gap worth watching across more sessions before redesigning the flow.

20.3 The “faster mode” question — worked through, not yet built

The most important thing to think through before building anything here. The instinct to stay conflicted about it was the right instinct.

- What's actually being described: someone making many high-stakes decisions a day doesn't need a deep council — they need a triage tool. That's a different product with a different posture. Quorum's value is depth, not speed, and a fast mode risks diluting exactly what makes it worth the price.
- What's worth preserving: the real signal isn't “make it faster” — it's “the text is too long to read in a high-pressure moment.” That's solvable without a mode change: shorter advisor outputs, text-to-speech, a summary card up top. Speed of reading isn't speed of thinking.
- The “easy / moderate / fast, like a game” framing is itself a signal — it means the person is thinking about this as a feature, not a core behaviour. A serious decision-support product and a game are not the same thing.

Decision: don't build a fast/game mode yet. Build shorter advisor output as a default first and see whether that resolves the “too long to read” complaint without introducing a mode. Revisit only after several more sessions show whether speed or readability is the actual issue.

20.4 Bugs & UX — logged

- Magic link for email linking failed during the session — reportedly fixed since; confirm end-to-end with a new user in the next session.
- Re-analysing a decision creates duplicate entries on the home page — the delete-button fix above addresses this; also clarify in the UX copy that re-analysing creates a new version, not the same decision.
- No back button after saving a decision record on the council page — users land in a dead end (re-analyse or new decision only). Low-effort fix.

20.5 Immediate next steps taken

- Logged the trade-offs insight and the delete-decision-button request into the product backlog before they got lost to the next session.
- Closed the fast-mode question with a decision (above) rather than leaving it open-ended.
- Scheduled a reconnect once she's run several more decisions through the product, rather than leaving it floating.